



---

# Welcome Email Sequence

— Alfocuss

[Project Deliverable Report](#)

Prepared: June 2026 • Version: 1.0 • Status: Complete

These are the five automated welcome emails new subscribers receive over their first two weeks after signing up to the Alfocuss newsletter. They are sent in the order shown below, beginning immediately on signup.

---

## Email 1 — Welcome — your AI shortlist is inside

---

Hi {{FIRST\_NAME}},

Welcome to Alfocuse. You signed up because you're trying to figure out which AI tools actually deserve a place in your academic workflow — and you've probably noticed that most "best AI tools" lists are written by people who haven't used the tools for serious research work.

We're trying to fix that. Every tool on Alfocuse has been tested across real academic tasks — literature reviews, manuscript drafts, interview transcription, peer review responses, the boring administrative emails that consume too much faculty time. We score each one across three dimensions: ease of use, academic value, and price-to-value. No tool gets a high score unless it earned one in actual use.

Three things to read first:

**1. Top Tools, ranked** The full ranked list across all four categories — writing, research, image generation, and coding assistants. If you only read one page on the site, read this one. → <https://aifocus.work/top-tools/>

**2. Our highest-scored tool right now: Elicit** Free tier covers most graduate-student use, structured extraction across academic papers, semantic search that finds work keyword searches miss. If you do literature reviews and have not tried it yet, it is the closest thing to a research assistant we have found on the AI side. → <https://aifocus.work/reviews/elicit-ai/>

**3. The 30-second sanity check** If you only have time for one read this week, the homepage covers the categories, the scoring methodology, and the top picks in each one. → <https://aifocus.work/>

What's coming over the next two weeks:

- **Wednesday** — the 30-minute literature review workflow that used to take a full day
- **Saturday** — the qualitative-research tool that transcribes a 60-minute interview in under 5 minutes
- **Next Tuesday** — the tools we actually use every day, and what each one replaced

If there's a specific tool, workflow, or question you want covered, just hit reply. Every email goes to a real person and we read all of them.

Thanks for being here.

Alan Breitler Founder, Alfocuse <https://aifocus.work>

— This email contains affiliate links. We may earn a small commission if you sign up through links in our emails or on the site, at no extra cost to you. We only recommend tools we genuinely use. Full disclosure: <https://aifocus.work/disclosure/>

## Email 2 — The 30-minute literature review workflow

---

Hi {{FIRST\_NAME}},

A confession: until late 2024, we treated AI tools for literature reviews as a curiosity rather than a serious option. The hallucinated citations were too frequent and the time spent fact-checking AI output usually exceeded the time it would have taken to do the search the old-fashioned way.

That has changed. Two tools in particular — Elicit and Consensus — have crossed the threshold where they save genuine time on real literature work, not theoretical time on demo searches. Here is the workflow we now use as the default first pass on any new research question.

### The 30-minute version

1. Start in Consensus (<https://aifocus.work/reviews/consensus-ai/>). Type your research question in plain English — not a Boolean keyword string. "Does mindfulness training reduce burnout in graduate students" works exactly as well as the cleaned-up version. Consensus returns the most relevant papers and, crucially, summarises whether the literature broadly supports, contests, or is mixed on your question. Five minutes.
2. Take the 5-10 most relevant Consensus results and paste them into Elicit (<https://aifocus.work/reviews/elicit-ai/>). Add columns for the data you actually care about — intervention type, sample size, study design, primary outcome, reported effect size. Elicit reads each paper and fills the columns. Ten to fifteen minutes for an initial table.
3. Spot-check Elicit's extractions for the top 3 papers. The tool is right about 90% of the time; the 10% where it is wrong matters. Five minutes.
4. You now have a structured table of the most relevant evidence on your question, in a fraction of the time it would have taken to read and extract manually. From here you decide whether to expand the search, refine the question, or start drafting.

We used this exact workflow last week on a research question we knew well, and it produced an evidence table comparable to one our team had previously spent two full days assembling by hand. Not a perfect replacement — there were a couple of methodological nuances the AI missed — but a 90%-accurate first pass in 30 minutes is genuinely transformative.

### Both tools have generous free tiers

This is the part most people miss. Consensus gives you 20 free searches per month. Elicit gives you 5,000 credits per month, enough for managing a medium-sized literature review. For occasional users — graduate students, faculty preparing a single grant proposal — the free tiers may be all you ever need.

If you do work that justifies a paid plan, Elicit Plus at \$10/month is the best-value tool we have reviewed. Full details and our scoring:

→ Elicit review: <https://aifocus.work/reviews/elicit-ai/> → Consensus review: <https://aifocus.work/reviews/consensus-ai/>

Try it on a question you already know the answer to. The workflow's value becomes obvious when you can immediately tell whether the AI got it right.

Next email (Saturday): the tool that transcribes a 60-minute research interview in under 5 minutes, plus when it works and when you absolutely should not use it.

If you have questions about adapting this workflow to your specific field, reply to this email.

Alan Breitler Founder, Aifocus <https://aifocus.work>

— This email contains affiliate links. We may earn a small commission if you sign up through these links, at no extra cost to you. We only recommend tools we genuinely use. Full disclosure: <https://aifocus.work/disclosure/>

---

## Email 3 — The interview transcript that wrote itself

---

Hi {{FIRST\_NAME}},

If you do qualitative research, this is the email in the sequence most likely to give back actual hours of your life.

The accepted rule of thumb for manual transcription is four to six hours per hour of audio. For a doctoral student running 30 interviews for a dissertation, that is a full month of work before the analysis even starts. Until recently, the only honest alternatives were either paying \$90+ per interview for human transcription (Rev.com and similar services) or living with the time cost.

That changed when Otter.ai became good enough.

### The honest version of how good it actually is

On clean audio with clear speakers, Otter produces transcripts at around 90 to 95% word accuracy. That is good enough to scan, code, and locate quotes. It is not good enough to publish a verbatim quote without checking. The workflow that works is:

1. Record your interview as you normally would.
2. Upload to Otter (or let it capture live via Zoom/Meet/Teams). Wait 3-5 minutes.
3. Skim the transcript at 1.5x audio playback, correcting errors as you go.
4. For a 60-minute interview, this takes about 25-35 minutes — roughly 4-6× faster than manual transcription from scratch.

What makes step 3 actually work is that the transcript is synchronised with the audio. Click any word and it jumps to that moment in the recording. We use this constantly during correction and later during coding — being able to click on a quote to verify it is the feature that turns Otter from "approximate text" into "research-usable transcript."

### Where it fails

Three failure modes you need to plan around:

- **Strong non-native accents.** Word accuracy drops to 75-85% for some accent profiles. Plan for higher correction time.
- **Disciplinary jargon.** Biomedical terminology, legal Latin, chemistry nomenclature — routinely misheard. We watched Otter transcribe "operationalisation" as "operational session" in a psychology interview. A small error, but it would have been embarrassing in a published quote.
- **Overlapping speech and focus groups.** Speaker identification confuses itself when three or more voices speak in quick succession. For focus group work, Otter saves time but does not eliminate the listening-back step.

### The free tier covers a lot more than people realise

Otter's free plan gives you 300 transcription minutes per month — enough for roughly 5 standard research interviews. If you only run interviews occasionally, the free tier may be all you need.

For active qualitative researchers, the Pro plan at \$8.33/month (billed annually) unlocks the audio file import feature, longer sessions, and advanced search. It is the cheapest credible AI transcription option we have found and it pays for itself in saved time after a single interview project.

→ Full review: <https://aifocus.work/reviews/otter-ai/>

### **One more thing — research ethics**

Otter is a third-party participant in your meetings. Your interview audio is processed on Otter's servers. Most universities permit this for low-sensitivity work, but research involving human subjects may require ethics-board approval before introducing a third-party transcription service. Check before you start.

For high-stakes work where every word must be exact (legal research, oral history projects, anything quoting vulnerable participants), Rev's human transcription at ~\$1.50/minute remains the safer choice. For routine research interviews, Otter wins on time and cost.

Next email (next Tuesday): the tools we actually reach for every day, and what each one replaced in our previous workflow.

Alan Breitler Founder, Aifocus <https://aifocus.work>

— This email contains affiliate links. We may earn a small commission if you sign up through these links, at no extra cost to you. We only recommend tools we genuinely use. Full disclosure: <https://aifocus.work/disclosure/>

---

## Email 4 — Tools we use every day (and what they replaced)

---

Hi {{FIRST\_NAME}},

The most useful question to ask about any AI tool is not "is it impressive" but "did it permanently change something in your workflow that used to work differently."

Here are the four tools that have cleared that bar for our team, what each one replaced, and the specific moments we reach for them.

**Claude Pro — drafting and analysis** Replaced: ChatGPT Plus as our default writing assistant.

We switched after running both side-by-side on academic writing tasks for three months. Claude produces noticeably fewer hallucinated citations, handles longer documents without losing context (200K token window), and the writing voice is closer to what we want for academic work — less chirpy, less prone to bullet-pointing everything. We still use ChatGPT for image-related work and a few specific use cases, but Claude is now the default.

Use it for: working through a difficult section of a draft, summarising a long document, asking "what are the methodological objections to this paper" before writing a peer review response.

→ <https://aifocus.work/reviews/claude-pro/>

**Grammarly — running in the background** Replaced: nothing. It is genuinely the only tool we have ever installed and forgotten about while still using daily.

Grammarly Premium catches the basic errors and the tone-register mismatches that even careful writers introduce in long drafts. The inline integration in Google Docs and Word means there is no workflow to learn — the tool just exists in the background and surfaces corrections as you write. For academics who write English as a second language, the Fluency suggestions in particular are excellent.

The one risk: it will sometimes try to "correct" discipline-specific terminology. Train its personal dictionary on your common terms and the false-positive rate drops sharply.

→ <https://aifocus.work/reviews/grammarly-ai/>

**QuillBot — focused revision sessions** Replaced: the ChatGPT-paste-and-paraphrase workflow we used before.

QuillBot is what we reach for when we need to revise a specific passage rather than improve overall writing. The seven paraphrasing modes — Standard, Fluency, Formal, Academic, Simple, Creative, Shorten — produce genuinely different outputs. The Shorten mode alone has rescued multiple abstracts trying to hit a 250-word limit.

The risk is real: Academic Mode is more conservative than the name suggests and can subtly change discipline-specific terminology. Treat every output as a draft requiring review, not a finished revision.

Free tier is enough for occasional use. The annual Premium plan at \$9.95/month effective is the realistic price point if you commit.

→ <https://aifocus.work/reviews/quillbot/>

**Perplexity Pro — the new search habit** Replaced: half of our Google Scholar usage.

For the early phase of any new research question — when you do not yet know what specific question to ask — Perplexity has become our default. Every answer includes inline citations to the actual sources, so you can verify what the tool is telling you. Pro plan lets you switch between GPT-4, Claude, Sonar, and Gemini per question, which is the underrated feature that justifies the upgrade.

We still use Google Scholar for known-author lookups and citation tracing, but for "I need to understand this topic enough to ask a real question" Perplexity is now the first stop.

→ <https://aifocus.work/reviews/perplexity-pro/>

### **The four together**

Used in sequence: Perplexity for orientation → Elicit for systematic extraction (covered in email 2) → Claude for synthesis and drafting → Grammarly + QuillBot for revision. This is the workflow that emerged unintentionally from using each tool for what it does best.

The financial picture: Grammarly Premium (\$12/mo) + Claude Pro (\$20/mo) + Perplexity Pro (\$20/mo) + QuillBot annual (\$9.95/mo effective) totals roughly \$62/month. For an active researcher producing regular written output, that is genuinely cheaper than the time-equivalent of doing this work without the tools — and many academics will get most of the value from just two or three of these subscriptions.

If you only adopt one new tool from this email, make it Claude Pro. The leap in quality over ChatGPT for academic writing is the biggest single upgrade we have experienced in two years.

Next email (Day 14): what's coming next on Alfocus, plus a small ask.

Alan Breitler Founder, Alfocus <https://aifocus.work>

— This email contains affiliate links. We may earn a small commission if you sign up through these links, at no extra cost to you. We only recommend tools we genuinely use. Full disclosure: <https://aifocus.work/disclosure/>

---

## Email 5 — What's coming next + a small ask

---

Hi {{FIRST\_NAME}},

You're at the end of the welcome sequence. From here, you'll get the weekly newsletter — one email every Thursday with new reviews, scoring updates, and a single piece of useful workflow advice.

Three things before we go.

### What's coming over the next month

We have new reviews in the pipeline for:

- An AI coding assistant comparison (Cursor vs Copilot vs Codeium) for academics who write analysis scripts
- A deep-dive on AI detection tools — what they actually catch, what they don't, and why universities are quietly dropping them
- A workflow guide for using Claude's Projects feature to build a persistent research workspace
- Updated scoring across the 11 tools currently reviewed, including any new releases and price changes

If any of these is more interesting to you than the others, hit reply and tell us — we move things up the queue based on subscriber requests more than anything else.

### The small ask

What's the AI tool, workflow, or comparison you most want covered that nothing on Alfocuse currently addresses? It could be a specific tool we haven't reviewed, a question that no review answers cleanly, or a workflow you've been trying to put together and can't find good guidance on.

Just hit reply. Every email goes to a real person and we read all of them. We write the most-requested topics first, so this is the single best way to influence what comes next.

### The honest expectations bit

A few things to know about how Alfocuse works:

- We make a small commission when someone signs up to a tool through our affiliate links. This is how the site is funded. It does not affect our scoring — tools that don't earn a high score don't get one, even when they pay us better. The full disclosure is at <https://aifocus.work/disclosure/>.
- We will never sell, share, or rent your email address. Newsletter signup is the only mailing list we run. Unsubscribe whenever you like.
- The weekly newsletter takes about 30 seconds to scan and 5 minutes to read fully if you want to dig into the linked content. If it ever stops being worth that 5 minutes, unsubscribe — we'd rather have a smaller list of people who actually read it.

Thanks for being here. The site is still relatively new and every subscriber matters more than you probably realise.

See you on Thursday.

Alan Breitler Founder, Alfocuse <https://aifocus.work>